

Differential Privacy and its Application in Aggregation

Part 2 — Privacy-preserving Aggregation

presenter: Le Chen

Nanyang Technological University

lechen0213@gmail.com

October 12, 2013

Outline

Introduction

Basic Construction

Distributed Differential Privacy

Discussion

Conclusion & Discussion

Reference



Cynthia Dwork.

Differential Privacy.

Invited talk at ICALP, Venice, Italy, July 10-14, 2006.

Automata, Languages and Programming, Lecture Notes in Computer Science Volume 4052, 2006, pp 1-12.



Wikipedia.

<http:

//en.wikipedia.org/wiki/Differential_privacy>



Elaine Shi, T-H. Hubert Chan, Eleanor Rieffel, Richard Chow and Dawn Song.

Privacy-Preserving Aggregation of Time-Series Data.

In Network and Distributed System Security Symposium (NDSS), 2011.

Reference




T-H. Hubert Chan, Elaine Shi, and Dawn Song

Privacy-Preserving Stream Aggregation with Fault Tolerance.

16th International Conference, FC 2012, Kralendijk, Bonaire, Februray 27-March 2, 2012. Financial Cryptography and Data Security, Lecture Notes in Computer Science Volume 7397, 2012, pp 200-214.

Motivation of Aggregation

- Statistics: In many practical applications, a data aggregator wishes to mine data coming from multiple organizations or individuals, to study **patterns** or **statistics** over a population.



Motivation of Aggregation

- ▶ Aggregator:
 - ▶ No aggregator.
 - ▶ Structure based aggregator.
 - ▶ Third-party aggregator.
- ▶ Advantage: Communication and computation overhead can be significantly decreased.
- ▶ Protect individual privacy:
 - ▶ Masking value.
 - ▶ Distributed differential privacy.

Two Techniques

- ▶ Basic construction - masking value.
- ▶ Distributed differential privacy.

Basic Construction - Masking Value

- ▶ A trusted dealer chooses a random generator $g \in \mathbb{G}$, and $n + 1$ random secrets $s_0, s_1, \dots, s_n \in \mathbb{Z}_p$ such that $s_0 + s_1 + s_2 + \dots + s_n = 0$.
- ▶ The aggregator obtains the capability $sk_0 := s_0$, and participant i obtains the secret key $sk_i := s_i$.
- ▶ NoisyEnc: $c \leftarrow g^{\hat{x}} H(t)^{sk_i}$, where $\hat{x} = x + r \pmod p$.
- ▶ AggrDec: $V = H(t)^{s_0} \prod_{i=1}^n c_i = g^{\sum_{i=1}^n \hat{x}_i}$.

Basic Construction - Masking Value

- ▶ A trusted dealer chooses a random generator $g \in \mathbb{G}$, and $n + 1$ random secrets $s_0, s_1, \dots, s_n \in \mathbb{Z}_p$ such that $s_0 + s_1 + s_2 + \dots + s_n = 0$.
- ▶ The aggregator obtains the capability $sk_0 := s_0$, and participant i obtains the secret key $sk_i := s_i$.
- ▶ NoisyEnc: $c \leftarrow g^{\hat{x}} H(t)^{sk_i}$, where $\hat{x} = x + r \pmod p$.
- ▶ AggrDec: $V = H(t)^{s_0} \prod_{i=1}^n c_i = g^{\sum_{i=1}^n \hat{x}_i}$.

Basic Construction - Masking Value

- ▶ A trusted dealer chooses a random generator $g \in \mathbb{G}$, and $n + 1$ random secrets $s_0, s_1, \dots, s_n \in \mathbb{Z}_p$ such that $s_0 + s_1 + s_2 + \dots + s_n = 0$.
- ▶ The aggregator obtains the capability $sk_0 := s_0$, and participant i obtains the secret key $sk_i := s_i$.
- ▶ NoisyEnc: $c \leftarrow g^{\hat{x}} H(t)^{sk_i}$, where $\hat{x} = x + r \pmod p$.
- ▶ AggrDec: $V = H(t)^{s_0} \prod_{i=1}^n c_i = g^{\sum_{i=1}^n \hat{x}_i}$.

Basic Construction - Masking Value

- ▶ A trusted dealer chooses a random generator $g \in \mathbb{G}$, and $n + 1$ random secrets $s_0, s_1, \dots, s_n \in \mathbb{Z}_p$ such that $s_0 + s_1 + s_2 + \dots + s_n = 0$.
- ▶ The aggregator obtains the capability $sk_0 := s_0$, and participant i obtains the secret key $sk_i := s_i$.
- ▶ NoisyEnc: $c \leftarrow g^{\hat{x}} H(t)^{sk_i}$, where $\hat{x} = x + r \pmod p$.
- ▶ AggrDec: $V = H(t)^{s_0} \prod_{i=1}^n c_i = g^{\sum_{i=1}^n \hat{x}_i}$.

Basic Construction - Masking Value

- ▶ To decrypt the sum $\sum_{i=1}^n \hat{x}_i$, it suffices to compute the discrete log of V base g .
- ▶ When the plaintext **space is small**, decryption can be achieved through a brute-force search.
- ▶ A better approach is to use Pollard's lambda method which requires decryption time roughly square root in the plaintext space.

Basic Construction - Masking Value

- ▶ To decrypt the sum $\sum_{i=1}^n \hat{x}_i$, it suffices to compute the discrete log of V base g .
- ▶ When the plaintext **space is small**, decryption can be achieved through a brute-force search.
- ▶ A better approach is to use Pollard's lambda method which requires decryption time roughly square root in the plaintext space.

Basic Construction - Masking Value


- ▶ To decrypt the sum $\sum_{i=1}^n \hat{x}_i$, it suffices to compute the discrete log of V base g .
- ▶ When the plaintext **space is small**, decryption can be achieved through a brute-force search.
- ▶ A better approach is to use Pollard's lambda method which requires decryption time roughly square root in the plaintext space.

Differential Privacy — Review

- ▶ ϵ -differential privacy ($\epsilon = \frac{\Delta(f)}{\lambda}$):

$$\Pr[\mathcal{A}(D_1) \in S] \leq e^\epsilon \times \Pr[\mathcal{A}(D_2) \in S],$$

- ▶ Laplace noise:



Motivation

- ▶ In previous differential privacy literature, a trusted aggregator is responsible for adding an **appropriate magnitude** of noise before releasing the statistics.
- ▶ Our approach is to let the participants add noise to their data before encrypting them (distributed).
- ▶ One naive solution is to rely on **a single participant** to add an appropriate magnitude of noise r to her data before submission.

Motivation

- ▶ In previous differential privacy literature, a trusted aggregator is responsible for adding an **appropriate magnitude** of noise before releasing the statistics.
- ▶ Our approach is to let the participants add noise to their data before encrypting them (distributed).
- ▶ One naive solution is to rely on **a single participant** to add an appropriate magnitude of noise r to her data before submission.

Motivation

- ▶ In previous differential privacy literature, a trusted aggregator is responsible for adding an **appropriate magnitude** of noise before releasing the statistics.
- ▶ Our approach is to let the participants add noise to their data before encrypting them (distributed).
- ▶ One naive solution is to rely on **a single participant** to add an appropriate magnitude of noise r to her data before submission.

Compromised Participants

- ▶ In particular, a subset of the participants may be compromised and collude with the data aggregator.
- ▶ In the **worst case**, if every participant believes that the other $n - 1$ participants may be compromised and collude with the aggregator, each participant would need to add **sufficient noise** to ensure the privacy of her own data.

Compromised Participants

- ▶ In particular, a subset of the participants may be compromised and collude with the data aggregator.
- ▶ In the **worst case**, if every participant believes that the other $n - 1$ participants may be compromised and collude with the aggregator, each participant would need to add **sufficient noise** to ensure the privacy of her own data.

Compromised Participants

- ▶ If at least γ fraction of the participants are honest and not compromised, then we can **distribute** the noise generation task amongst these participants. Each participant may add less noise, and as long as the noise in the final statistic is large enough, individual privacy is protected.
- ▶ Our scheme assumes that the participants have an **a priori estimate** on the lower bound for γ .

Compromised Participants

- ▶ If at least γ fraction of the participants are honest and not compromised, then we can **distribute** the noise generation task amongst these participants. Each participant may add less noise, and as long as the noise in the final statistic is large enough, individual privacy is protected.
- ▶ Our scheme assumes that the participants have an **a priori estimate** on the lower bound for γ .

Algebraic Constraints

- ▶ Most encryption schemes require that the plaintext be picked from a group comprised of **discrete elements**.
- ▶ We choose to use a **symmetric geometric distribution** instead of the more commonly used Laplace distribution.

Algebraic Constraints

- ▶ Most encryption schemes require that the plaintext be picked from a group comprised of **discrete elements**.
- ▶ We choose to use a **symmetric geometric distribution** instead of the more commonly used Laplace distribution.

Symmetric Geometric Distribution

Definition:

- ▶ Let $\alpha > 1$. We denote by $\text{Geom}(x, \alpha)$ the symmetric geometric distribution that takes integer values x such that the probability mass function at x is $\frac{\alpha-1}{\alpha+1} \cdot \alpha^{-|x|}$.
- ▶ The symmetric geometric distribution $\text{Geom}(x, \alpha)$ can be viewed as a discrete version of the Laplace distribution $\text{Lap}(x, \lambda)$ (where $\alpha \approx \exp(\frac{1}{\lambda})$), whose probability density function is $x \mapsto \frac{1}{2\lambda} \exp(-\frac{|x|}{\lambda})$.

Symmetric Geometric Distribution


Definition:

- ▶ Let $\alpha > 1$. We denote by $\text{Geom}(x, \alpha)$ the symmetric geometric distribution that takes integer values x such that the probability mass function at x is $\frac{\alpha-1}{\alpha+1} \cdot \alpha^{-|x|}$.
- ▶ The symmetric geometric distribution $\text{Geom}(x, \alpha)$ can be viewed as a discrete version of the Laplace distribution $\text{Lap}(x, \lambda)$ (where $\alpha \approx \exp(\frac{1}{\lambda})$), whose probability density function is $x \mapsto \frac{1}{2\lambda} \exp(-\frac{|x|}{\lambda})$.

Differential Privacy

- ▶ Let $\epsilon > 0$. Suppose u and v are two integers such that $|u - v| \leq \Delta$. Let r be a random variable having distribution $\text{Geom}(\alpha)$, where $\alpha \approx \exp(\frac{1}{\lambda}) = \exp(\frac{\epsilon}{\Delta})$. Then, for any integer k , $\Pr[u + r = k] \leq \exp(\epsilon) \cdot \Pr[v + r = k]$.

Laplace Distribution



Laplace Distribution Probability Density Function

Error

- ▶ Our mechanism ensures small error of roughly $O(\frac{\Delta}{\epsilon} \sqrt{\frac{1}{\gamma}})$ magnitude.
- ▶ Consider the extreme case when $\gamma = O(\frac{1}{n})$, i.e., each participant believes that all other participants may be compromised. Then, our accumulated noise would be $O(\frac{\Delta}{\epsilon} \sqrt{\frac{1}{\gamma}}) = O(\frac{\Delta}{\epsilon} \sqrt{n})$.
- ▶ According to the **central limit theorem**, the sum of n independent symmetric noises of magnitude $O(\frac{\Delta}{\epsilon})$ results in a final noise of magnitude $O(\frac{\Delta}{\epsilon} \sqrt{n})$ with high probability.

Error

- ▶ Our mechanism ensures small error of roughly $O(\frac{\Delta}{\epsilon} \sqrt{\frac{1}{\gamma}})$ magnitude.
- ▶ Consider the extreme case when $\gamma = O(\frac{1}{n})$, i.e., each participant believes that all other participants may be compromised. Then, our accumulated noise would be $O(\frac{\Delta}{\epsilon} \sqrt{\frac{1}{\gamma}}) = O(\frac{\Delta}{\epsilon} \sqrt{n})$.
- ▶ According to the **central limit theorem**, the sum of n independent symmetric noises of magnitude $O(\frac{\Delta}{\epsilon})$ results in a final noise of magnitude $O(\frac{\Delta}{\epsilon} \sqrt{n})$ with high probability.

Error

- ▶ Our mechanism ensures small error of roughly $O(\frac{\Delta}{\epsilon} \sqrt{\frac{1}{\gamma}})$ magnitude.
- ▶ Consider the extreme case when $\gamma = O(\frac{1}{n})$, i.e., each participant believes that all other participants may be compromised. Then, our accumulated noise would be $O(\frac{\Delta}{\epsilon} \sqrt{\frac{1}{\gamma}}) = O(\frac{\Delta}{\epsilon} \sqrt{n})$.
- ▶ According to the **central limit theorem**, the sum of n independent symmetric noises of magnitude $O(\frac{\Delta}{\epsilon})$ results in a final noise of magnitude $O(\frac{\Delta}{\epsilon} \sqrt{n})$ with high probability.

Distributed Differential Privacy

Definition $((\epsilon, \delta)$ -DD-Privacy)

- ▶ Suppose $\epsilon > 0$, $0 \leq \delta < 1$ and $0 < \gamma \leq 1$. We say the function f achieves (ϵ, δ) -distributed differential privacy (DD-privacy) under γ fraction of uncompromised participants if the following condition holds.

$$\Pr[f(\hat{x}) \in S] \leq \exp(\epsilon) \cdot \Pr[f(\hat{y}) \in S] + \delta.$$

Distributed Differential Privacy

Algorithm 1: DD-Private Data Randomization Procedure.

Let $\alpha := \exp(\frac{\epsilon}{\Delta})$ and $\beta := \frac{1}{\gamma n} \log \frac{1}{\delta}$.

Let $\mathbf{x} = (x_1, \dots, x_n)$ denote all participants' data in a certain time period.

foreach participant $i \in [n]$ **do**

 Sample noise r_i according to the following distribution.

$$r_i \leftarrow \begin{cases} \text{Geom}(\alpha) & \text{with probability } \beta \\ 0 & \text{with probability } 1 - \beta \end{cases}$$

 Randomize data by computing $\hat{x}_i \leftarrow x_i + r_i \pmod p$.

- ▶ **Lemma.** Let $\epsilon > 0$ and $0 < \delta < 1$. Suppose at least γ fraction of participants are uncompromised. Then, the above randomization procedure achieves (ϵ, δ) -DD-privacy with respect to **sum**, for $\beta = \min\{\frac{1}{\gamma n} \log \frac{1}{\delta}, 1\}$

Parameters

- ▶ $\epsilon > 0$.
- ▶ $0 < \gamma \leq 1$.
- ▶ $0 < \delta < 1$.
- ▶ $\beta = \frac{1}{\gamma n} \log \frac{1}{\delta}$.
- ▶ Δ .
- ▶ $\alpha \approx \exp(\frac{1}{\lambda}) = \exp(\frac{\epsilon}{\Delta})$.

The Parameter ϵ

- ▶ $\epsilon > 0$.
- ▶ The privacy parameter.

The Parameter γ

- ▶ $0 < \gamma \leq 1$.
- ▶ The proportion of trusted participants.

The Parameter δ

- ▶ $0 < \delta < 1$.
- ▶ Another privacy parameter.

The Parameter β

- ▶ $\beta = \frac{1}{\gamma n} \log \frac{1}{\delta}.$
- ▶ The probability that a trusted participant generates a noise.

Relationship of γ , δ , and β

- ▶ $\beta = \frac{1}{\gamma n} \log \frac{1}{\delta}.$
- ▶ Given γ .
- ▶ Given δ .

The Parameter Δ

- ▶ $\Delta = \max\{|f(\hat{x}) - f(\hat{y})|\}.$
- ▶ The probability that a trusted participant generates a noise.

The Parameter α

- ▶ $\alpha \approx \exp(\frac{1}{\lambda}) = \exp(\frac{\epsilon}{\Delta})$.
- ▶ The magnitude of noise, the larger α is the smaller the noise is.

Relationship of α , and Δ

- ▶ $\alpha \approx \exp(\frac{1}{\lambda}) = \exp(\frac{\epsilon}{\Delta})$.
- ▶ The larger Δ is, the larger the noise is needed.

Confliction

- ▶ Can the basic construction extend to support distributed differential privacy?
- ▶ The basic construction needs all users to participate, or the masking value cannot be canceled.
- ▶ So it is impossible for the basic construction to support distributed differential privacy.

Confliction

- ▶ Can the basic construction extend to support distributed differential privacy?
- ▶ The basic construction needs all users to participate, or the masking value cannot be canceled.
- ▶ So it is impossible for the basic construction to support distributed differential privacy.

Confliction

- ▶ Can the basic construction extend to support distributed differential privacy?
- ▶ The basic construction needs all users to participate, or the masking value cannot be canceled.
- ▶ So it is impossible for the basic construction to support distributed differential privacy.

Conclusion & Discussion

- ▶ We introduced the basic construction that uses masking value and the distributed differential privacy.
- ▶ Achieving distributed differential privacy with small error is not easy.
- ▶ It also depends on the query situation.
- ▶ Discussion?

Conclusion & Discussion

- ▶ We introduced the basic construction that uses masking value and the distributed differential privacy.
- ▶ Achieving distributed differential privacy with small error is not easy.
- ▶ It also depends on the query situation.
- ▶ Discussion?

Conclusion & Discussion

- ▶ We introduced the basic construction that uses masking value and the distributed differential privacy.
- ▶ Achieving distributed differential privacy with small error is not easy.
- ▶ It also depends on the query situation.
- ▶ Discussion?

Conclusion & Discussion

- ▶ We introduced the basic construction that uses masking value and the distributed differential privacy.
- ▶ Achieving distributed differential privacy with small error is not easy.
- ▶ It also depends on the query situation.
- ▶ Discussion?