

Differentially Private Aggregation of Distributed Time-series with Transformation and Encryption

Part 4 — Differentially Private Aggregation 2

presenter: Le Chen

Nanyang Technological University

lechen0213@gmail.com

October 27, 2013

Outline

Introduction

L_2 Sensitivity

Discrete Fourier Transform

Conclusion & Discussion

Reference



Cynthia Dwork.

Differential Privacy.

Invited talk at ICALP, Venice, Italy, July 10-14, 2006.

Automata, Languages and Programming, Lecture Notes in Computer Science Volume 4052, 2006, pp 1-12.



Wikipedia.

<[http:](http://en.wikipedia.org/wiki/Differential_privacy)

[//en.wikipedia.org/wiki/Differential_privacy](http://en.wikipedia.org/wiki/Differential_privacy)>



Elaine Shi, T-H. Hubert Chan, Eleanor Rieffel, Richard Chow and Dawn Song.

Privacy-Preserving Aggregation of Time-Series Data.

In Network and Distributed System Security Symposium (NDSS), 2011.

Reference



T-H. Hubert Chan, Elaine Shi, and Dawn Song

Privacy-Preserving Stream Aggregation with Fault Tolerance.

16th International Conference, FC 2012, Kralendijk, Bonaire, Februray 27-March 2, 2012. Financial Cryptography and Data Security, Lecture Notes in Computer Science Volume 7397, 2012, pp 200-214.



Vibhor Rastogi, Suman Nath

Differentially Private Aggregation of Distributed Time-series with Transformation and Encryption.

ACM SIGMOD 2010, New York, NY. Proceedings of the 2010 ACM SIGMOD International Conference on Management of Data, pages 735-746.

Motivation of Aggregation




Figure 1: System Model (Users with data I_1, \dots, I_U Aggregator issues recurring query $Q = Q_1, \dots, Q_n$ No trusted server has $I = I_1 \cup I_2 \dots \cup I_U$ to evaluate $Q(I)$)

Figure: System Model

Notations

- ▶ $I = I_1 \cup I_2 \cdots \cup I_U$ is a set of vectors.
- ▶ $I_j = [I_{j1}, I_{j2}, \cdots, I_{jn}]$ is a $n \times 1$ vector, where $I_{ij} \in \{0, 1\}$ represents whether the average weight of user is over 200 lb in month j .
- ▶ $nbrs(I)$: the data set obtained from adding/removing one user's data from I .
- ▶ $Q_j(I)$ outputs "the number of users over 200 lb" in month j , i.e. $\sum_{n=1}^U I_{ikj}$.
- ▶ $Q(I) = \{Q_1(I), Q_2(I), \cdots Q_n(I)\}$.

LPA (Laplace Perturbation Algorithm)

- ▶ By adding $LAP(\lambda)$ noise to the final results, $LPA(Q, \lambda)$ is ϵ -differentially private for $\lambda = \Delta(Q)/\epsilon$.
- ▶ $\tilde{Q} = Q(I) + LAP^n(\lambda)$.
- ▶ DLPA: Each user u adds its noise n_u to x_u .
- ▶ $\tilde{Q} = \sum_u (x_u + n_u) = Q(I) + \sum_u n_u$.

Basic Distributed Protocol

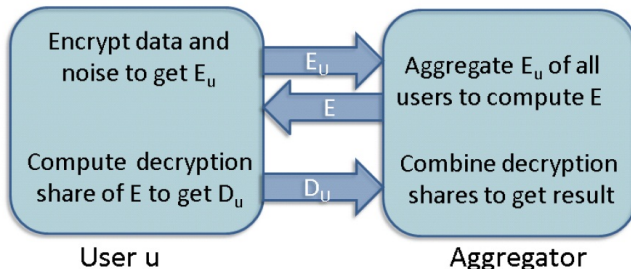


Figure 2: Basic Distributed Protocol (homomorphic property exploited to aggregate users' encryption & threshold property to combine users' decryption shares)

Distributed Decryption

- ▶ Distributed keys: The private key λ is shared by U users as $\lambda = \sum_u \lambda_u$.
- ▶ Each user u computes his decryption share $c_u = c^{\lambda_u}$.
- ▶ The aggregator combines $c' = \prod_{u=1}^U c_u$.
- ▶ The final result $t = \frac{L(c' \bmod N^2)}{L(g^\lambda \bmod N^2)} \bmod N$.

Protocol for Computing Exact Sum

- ▶ $\text{Encrypt-Sum}(x_u, r_u)$: the encryption of $\sum_{u=1}^U (x_u + r_u) = Q + \sum_{u=1}^U r_u$.
- ▶ $\text{Decrypt-Sum}(c, r_u)$: each user u computes a decryption share $c'_u = c^{\lambda_u} g^{-r_u \lambda}$.
- ▶ Final result: the aggregator aggregates $c' = \prod_{u=1}^U c'_u$ and gets $Q = \frac{L(c' \bmod N^2)}{L(g^\lambda \bmod N^2)} \bmod N$.

Protocol for Computing Noisy Sum

- ▶ Let $Y_i \sim N(0, \lambda)$ for $i \in \{1, 2, 3, 4\}$ be four Gaussian random variables. Then $Z = Y_1^2 + Y_2^2 - Y_3^2 - Y_4^2$ is a $Lap(2\lambda^2)$ random variable.

Algorithm 5.4 Encrypt-Noisy-Sum(x_u, r_u)

- 1: User u chooses five random numbers $r_u^1, r_u^2, \dots, r_u^5$ from \mathbb{Z}_m and computes $r_u = r_u^1 + r_u^2 - r_u^3 - r_u^4 + r_u^5$.
 - 2: User u generates four $N(0, \sqrt{2\lambda}/U)$ random variables y_u^1, \dots, y_u^4 .
 - 3: Let $c^j = \text{Encrypt-Sum-Squared}(y_u^j, r_u^j)$ for $j \in \{1, 2, 3, 4\}$.
 - 4: Let $c^5 = \text{Encrypt-Sum}(x_u, r_u^5)$.
 - 5: Aggregator computes $c = \frac{c^1 c^2 c^5}{c^3 c^4}$.
-

Protocol for Computing Noisy Sum

Algorithm 5.3 Encrypt-Sum-Squared(y_u, r_u) Protocol

- 1: User u computes $c_u = Enc(y_u + a_u + b_u)$ and sends it to the aggregator.
 - 2: The aggregator computes $c = \prod_{u=1}^U c_u$ and sends it to each user u .
 - 3: Each user u generates a random $r_u \in \mathbb{Z}_m$, computes $c_u = c^{y_u - a_u + b_u} Enc(r_u)$.
 - 4: The aggregator collects c_u from each user and computes $c' = (\prod_{u=1}^U c_u) Enc(a^2)$
-

where $a = \sum_u a_u$, $Enc(a^2)$ is public, and $\sum_u b_u = 0$.

L_2 Sensitivity

- ▶ If Q is a query sequence, $Q(I)$ and $Q(I')$ are each vectors.
- ▶ L_1 distance metric, denoted as $|Q(I) - Q(I')|_1$, that measures the Manhattan distance $\sum_j |Q_j(I) - Q_j(I')|$ between these vectors.
- ▶ L_2 distance metric, denoted as $|Q(I) - Q(I')|_2$, that measures the Euclidean distance $\sqrt{\sum_j (Q_j(I) - Q_j(I'))^2}$.

Sensitivity

DEFINITION 2.2 (SENSITIVITY [7]). Let \mathbf{Q} be any query sequence. For $p \in \{1, 2\}$, the L_p sensitivity of \mathbf{Q} , denoted $\Delta_p(\mathbf{Q})$, is the smallest number such that for all I and $I' \in \text{nbrs}(I)$,

$$\|\mathbf{Q}(I) - \mathbf{Q}(I')\|_p \leq \Delta_p(\mathbf{Q})$$

For a single snapshot query Q_i , the L_1 and L_2 sensitivities are the same, and we write $\Delta(Q_i) = \Delta_1(Q_i) = \Delta_2(Q_i)$.

Sensitivity Example

EXAMPLE 2.1. Consider a query Q counting the number of users whose weight in month 1 is greater than 200 lb. Then $\Delta(Q)$ is simply 1 as Q can differ by at most 1 on adding/removing a single user's data. Now consider $\mathbf{Q} = Q_1, \dots, Q_n$, where Q_i counts users whose weight in month i is greater than 200 lb. Then $\Delta_1(\mathbf{Q})$ is n (for the pair I, I' which differ in a single user having weight > 200 in each month i) and $\Delta_2(\mathbf{Q}) = \sqrt{n}$ (for the same pair I, I').

DFT (Discrete Fourier Transform)

- ▶ The DFT of an n -dimensional sequence X is a linear transform giving another n -dimensional sequence.
- ▶ $DFT(X)$: the j^{th} element $DFT(X)_j = \sum_{i=1}^n e^{\frac{2\pi\sqrt{-1}}{n}ji} X_i$.
- ▶ $DFT(X)$: the j^{th} element of the Inverse DFT is $IDFT(X)_j = \frac{1}{n} \sum_{i=1}^n e^{-\frac{2\pi\sqrt{-1}}{n}ji} X_i$.
- ▶ Furthermore, $IDFT(DFT(X)) = X$.

$DFT^k(X)$

- ▶ Denote $DFT^k(X)$ as the first k elements of $DFT(X)$.
- ▶ The elements of $DFT^k(X)$ are called the Fourier coefficients of the k **lowest frequencies** and they compactly represent the **high-level trends** in X .
- ▶ An approximation X' to X can be obtained from $DFT^k(X)$ as follows:
 - ▶ $PAD^n(DFT^k(X))$: appends $n - k$ 0 to $DFT^k(X)$,
 - ▶ compute $X' = IDFT(PAD^n(DFT^k(X)))$.


$DFT^k(X)$

- ▶ Obviously X' may be different from X as ignoring the last $n - k$ Fourier coefficients may introduce some error.
- ▶ Denote $RE_j^k(X)$, short for reconstruction error at the j^{th} position, to be the value $|X'_j - X_j|$.

Reconstruction Error

EXAMPLE 4.1. *To give a sense of the reconstruction error, we consider a sequence \mathbf{X} of length $n = 2000$ representing the number of people with weight > 200 in a real dataset (more details in Section 7), counted once every day over 2000 days. Fig. 3 shows the reconstructed sequence, \mathbf{X}' , using $k = 20$ DFT coefficients along with the original sequence \mathbf{X} . \mathbf{X} shows the temporal trend in the # of overweight people in the dataset. As shown, \mathbf{X}' captures the trend accurately showing that the reconstruction error is small even when compressing from $n = 2000$ to $k = 20$ DFT coefficients.*

Reconstruction



FPA (Fourier Perturbation Algorithm)

Algorithm 4.1 FPA_k (Inputs: sequence \mathbf{Q} , parameter λ)

- 1: Compute $\mathbf{F}^k = \mathbf{DFT}^k(\mathbf{Q}(I))$.
 - 2: Compute $\tilde{\mathbf{F}}^k = \text{LPA}(\mathbf{F}^k, \lambda)$
 - 3: Return $\tilde{\mathbf{Q}} = \mathbf{IDFT}(\text{PAD}^n(\tilde{\mathbf{F}}^k))$
-

- ▶ The parameter λ in FPA_k needs to be adjusted in order to get ϵ -differential privacy.
- ▶ Since FPA_k perturbs the sequence F^k , λ has to be calibrated according to the L_1 sensitivity, $\Delta_1(F^k)$, of F^k .

THEOREM 4.1. Denote $\mathbf{F}^k = \mathbf{DFT}^k(\mathbf{Q}(I))$ the first k DFT coefficients of $\mathbf{Q}(I)$. Then, (i) the L_1 sensitivity, $\Delta_1(\mathbf{F}^k)$, is at most \sqrt{k} times the L_2 sensitivity, $\Delta_2(\mathbf{Q})$, of \mathbf{Q} , and (ii) $FPA_k(\mathbf{Q}, \lambda)$ is ϵ -differentially private for $\lambda = \sqrt{k}\Delta_2(\mathbf{Q})/\epsilon$.

Proof

- ▶ (i) holds since $\Delta_2(F^k) \leq \Delta_2(Q)$, as the n Fourier coefficients have the same L_2 norm as Q , while F^k ignores the last $n - k$ Fourier coefficients,
- ▶ and $\Delta_1(F^k) \leq \sqrt{k}\Delta_2(F^k)$, due to a standard inequality between the L_1 and L_2 norms of a sequence.
- ▶ Let $x_j = |Q_j(I) - Q_j(I')|$, then $\Delta_1(Q) = \sum_j |Q_j(I) - Q_j(I')| = \sum_j x_j$, $\Delta_2(Q) = \sqrt{\sum_j (Q_j(I) - Q_j(I'))^2} = \sqrt{\sum_j x_j^2}$. Since the inequality

$$\frac{\sum_j x_j}{n} \leq \sqrt{\frac{\sum_j x_j^2}{n}}$$

holds, we can see that $\Delta_1(F^k) \leq \sqrt{k}\Delta_2(F^k)$.

Proof

- ▶ (ii) follows since for $\lambda = \sqrt{k}\Delta_2(Q)/\epsilon \geq \Delta_1(F^k)/\epsilon$,
- ▶ $\tilde{F}^k = LPA(F^k, \lambda)$ computed in step 2 is ϵ -differential private, and \tilde{Q} in step 3 is obtained using \tilde{F}^k only.

Accuracy

THEOREM 4.2. Fix $\lambda = \sqrt{k}\Delta_2(\mathbf{Q})/\epsilon$ so that $FPA_k(\mathbf{Q}, \lambda)$ is ϵ -differentially private. Then for all $i \in \{1, \dots, n\}$, the error $i(FPA_k)$ is $k/\epsilon + RE_i^k(\mathbf{Q}(I))$.

- ▶ The theorem shows that the error by FPA_k for each query is $k/\epsilon + RE_i^k(Q(I))$, but the LPA yields an error of n/ϵ .
- ▶ Since the reconstruction error, $RE_i^k(Q(I))$, is often small even for $k \ll n$, we expect the error in FPA_k to be much smaller than in LPA.
- ▶ This hypothesis is confirmed in our experiments that show that FPA_k gives **orders of magnitude improvement** over LPA in terms of error.

Choosing the Right k

- ▶ So far we have assumed that k is known to us.
- ▶ Since $error_i(FPA_k)$ is $k/\epsilon + RE_i^k(Q(I))$, a good value of k is important in obtaining a good trade-off between the perturbation error, k/ϵ , and the reconstruction error, $RE_i^k(Q(I))$.
- ▶ If k is too big, the perturbation error becomes too big, while if k is too small the reconstruction error becomes too high.

Choosing the Right k

- ▶ We can often choose k based on prior assumptions about $Q(I)$.
- ▶ For instance, if $Q(I)$ is such that the Fourier coefficients corresponding to $Q(I)$ **decrease exponentially fast**, then only a constant number (say $k = 10$) of Fourier coefficients need to be retained during perturbation.
- ▶ Our experiments show that this naive method is applicable in many practical scenarios as Fourier coefficients of many real-word sequences decrease very rapidly.

Conclusion & Discussion

- ▶ Last week, we introduced an aggregation protocol supports distributed differential privacy and distributed decryption.
- ▶ This week we talked about using (DFT) Discrete Fourier Transform in LPA (Laplace Perturbation Algorithm) to get FPA (Fourier Perturbation Algorithm), and discussed the error of DFT and chosen of k .
- ▶ FPA can achieve the same differential privacy level with less error, the noise actually reduces by a factor of n/k .